

GKR

Chap 4 Interactive Proofs

Pushpendra Pal

June 2025

Reading notes from Justin Thaler's "Proofs Argument and Zero Knowledge"

Keywords: Interactive Proofs, Wiring predicate, Multiple polynomial evaluation reduction, Multi-Linear Extension, Sum-check protocol

1 Prerequisites

- Reducing multiple polynomial evaluation to one
- Sum check protocol
- Multi-linear extension

2 Introduction

2.1 Objective

The verifier needs to know that a circuit has been evaluated correctly, without him actually performing the circuit evaluation. The verifier does know the circuit and the input values but is lazy enough to not do the evaluation himself.

2.2 Idea

- The prover takes a layered circuit in which there is only connection between the adjacent layers, and claims the values at each level one by one in each round.
- The circuit is known to both the verifier and the prover.
- The values at each level are encoded in an MLE of all the values with their position (W_i , where i is the level starting from the top 0). There could be multiple outputs at the top level also. The last level contains the inputs, so the Prover in the last round would send the MLE of all the inputs.

•

$$\widetilde{W}_i(r_i) = f_{r_i}^i = \widetilde{add}_i(r_i, b, c) \cdot [\widetilde{W}_{i+1}(b) + \widetilde{W}_{i+1}(c)] + \widetilde{mult}_i(r_i, b, c) \cdot [\widetilde{W}_{i+1}(b) \cdot \widetilde{W}_{i+1}(c)] \quad (1)$$

3 Wiring predicate

The gates are also encoded as functions of three variables, r_i, b, c . r_i is the **output position** and b, c are the input **positions, not the input** to the specific gate. r_i belongs to the higher layer (closer to the root layer), while b, c belong to the lower layer. The positions are typically encoded as binary.

$add(r_i, b, c) = 1$ if wires are connected and 0 otherwise.




Figure 4.12: Example circuit $\mathcal{C}$ and input x , and resulting functions W_i for each layer i of $\mathcal{C}$. Note that $\mathcal{C}$ has two output gates.

Similarly, mult_1 is a function on domain $\{0,1\}^2 \times \{0,1\}^2 \times \{0,1\}^2$. It evaluates to 0 on all inputs except for the following four, on which it evaluates to 1:

- $((0,0), (0,0), (0,0))$.
- $((0,1), (0,1), (0,1))$.
- $((1,0), (0,1), (1,0))$.
- $((1,1), (1,1), (1,1))$.

Note that for each layer i , add_i and mult_i depend only on the circuit $\mathcal{C}$ and not on the input x to $\mathcal{C}$. In contrast, the function W_i *does* depend on x . This is because W_i maps each gate label at layer i to the value of the gate when $\mathcal{C}$ is evaluated on input x .

4 Protocol

Description of the GKR protocol, when applied to a layered arithmetic circuit C of depth d and fan-in two on input $x \in \mathbb{F}^n$. Throughout, k_i denotes $\log_2(S_i)$ where S_i is the number of gates at layer i of C .

- At the start of the protocol, $\mathcal{P}$ sends a function $D: \{0, 1\}^{k_0} \rightarrow \mathbb{F}$ claimed to equal $\tilde{W}_0$ (the function mapping output gate labels to output values). $\mathcal{P}$ at $k_0=0$ $\{0, 1\}^0 = \{\emptyset\}$, $\mathcal{P}$ just sends the o/p value.
- $\mathcal{V}$ picks a random $r_0 \in \mathbb{F}^{k_0}$ and lets $m_0 \leftarrow \tilde{D}(r_0)$. The remainder of the protocol is devoted to confirming that $m_0 = \tilde{W}_0(r_0)$. $\mathcal{V}$ doesn't send anything in case $k_0=0, S_1=1$.
- For $i=0, 1, \dots, d-1$:

- Define the $(2k_{i+1})$ -variate polynomial $f_i^{(i)}(b, c)$ is $\log_2(S_{i+1})$ bits, $\tilde{W}_{i+1}(b, c)$ is $2k_{i+1}$ bits

$$f_i^{(i)}(b, c) := \widetilde{\text{add}}_i(r_i, b, c) (\tilde{W}_{i+1}(b) + \tilde{W}_{i+1}(c)) + \widetilde{\text{mult}}_i(r_i, b, c) (\tilde{W}_{i+1}(b) \cdot \tilde{W}_{i+1}(c)).$$

- $\mathcal{P}$ claims that $\sum_{b, c \in \{0, 1\}^{k_{i+1}}} f_i^{(i)}(b, c) = m_i$. This is done to check.

Here $\sum$ -check could be $x(0, c) + x(1, c)$ So that $\mathcal{V}$ may check this claim, $\mathcal{P}$ and $\mathcal{V}$ apply the sum-check protocol to $f_i^{(i)}$, up until $\mathcal{V}$'s final check in that protocol, when $\mathcal{V}$ must evaluate $f_i^{(i)}$ at a randomly chosen point $(b^*, c^*) \in \mathbb{F}^{k_{i+1}} \times \mathbb{F}^{k_{i+1}}$. See Remark (a) at the end of this codebox. (answers do no)

- Let ℓ be the unique line satisfying $\ell(0) = b^*$ and $\ell(1) = c^*$. $\mathcal{P}$ sends a univariate polynomial q of degree at most k_{i+1} to $\mathcal{V}$, claimed to equal $\tilde{W}_{i+1}$ restricted to ℓ . (This is done to reduce D in degree of W_0)
- $\mathcal{V}$ now performs the final check in the sum-check protocol, using $q(0)$ and $q(1)$ in place of $\tilde{W}_{i+1}(b^*)$ and $\tilde{W}_{i+1}(c^*)$. See Remark (b) at the end of this codebox.

- $\mathcal{V}$ chooses $r^* \in \mathbb{F}$ at random and sets $r_{i+1} = \ell(r^*)$ and $m_{i+1} \leftarrow q(r_{i+1})$.


- $\mathcal{V}$ checks directly that $m_d = \tilde{W}_d(r_d)$ using Lemma 3.8.


Note that $\tilde{W}_d$ is simply $\tilde{x}$, the multilinear extension of the input x when x is interpreted as the evaluation table of a function mapping $\{0, 1\}^{\log n} \rightarrow \mathbb{F}$.

Remark a. Note that $\mathcal{V}$ does not actually know the polynomial $f_i^{(i)}$, because $\mathcal{V}$ does not know the polynomial $\tilde{W}_{i+1}$ that appears in the definition of $f_i^{(i)}$. However, the sum-check protocol does not require $\mathcal{V}$ to know anything about the polynomial to which it is being applied, until the very final check in the protocol (see Remark 4.2).

Remark b. We assume here that for each layer i of C , $\mathcal{V}$ can evaluate the multilinear extensions $\widetilde{\text{add}}_i$ and $\widetilde{\text{mult}}_i$ at the point (r_i, b^*, c^*) in polylogarithmic time. Hence, given $\tilde{W}_{i+1}(b^*)$ and $\tilde{W}_{i+1}(c^*)$, $\mathcal{V}$ can quickly evaluate $f_i^{(i)}(b^*, c^*)$ and thereby perform its final check in the sum-check protocol applied to $f_i^{(i)}$.

Figure 4.13: Self-contained description of the GKR protocol for arithmetic circuit evaluation.





- $\text{add}(a_i, b, c) = \begin{cases} 1 & \text{if } a_i = b+c \\ 0 & \text{otherwise} \end{cases}$
 $a_i = b+c$
 $b = w_1(a_i) \oplus c = w_2(a_i)$

- $f_{a_i}^{(1)}$ derived to be equal $w_i(a_i)$
- $f_{a_i}^{(1)}(b, c) = \text{add}(a_i, b, c) = \widetilde{w}_i(b) + \widetilde{w}_i(c)$
 $\text{mult}(a_i, b, c) = (\widetilde{w}_i(b) \cdot \widetilde{w}_i(c))$

- $\deg(f_{a_i}^{(1)}) = 2k_{i+1}$
- $(b^*, c^*) \in \mathbb{F}^{2k_{i+1}} \times \mathbb{F}^{k_{i+1}}$

- ℓ is unique line passing through b, c
 $\ell(0) = b^*, \ell(1) = c^*$
- P makes $q_i = \text{MOL}(\text{representative form of } \ell)$
 $\deg(q_i) = k_{i+1}$


- See example of making q_i from ℓ and w in text.

PROVER
 DEFINE $D: \{0, 1\}^d \rightarrow \mathbb{F}$
 (claim to equal $w(0)$)



DEFINE $f_{a_i}^{(1)}$
 (claim to equal $w_i(a_i)$)
 $\sum_{b, c \in \{0, 1\}^{k_{i+1}}} f_{a_i}^{(1)}(b, c) = m_i$

DEF $\ell(0) = b^*, \ell(1) = c^*$
 DEF $q_i(0) = \widetilde{w}_i(b^*)$
 $q_i(1) = \widetilde{w}_i(c^*)$



VERIFIER

$\widetilde{D} \xleftarrow{\text{MLE}} D$
 $q_i, 0 \leftarrow \mathbb{F}^{k_{i+1}}$
 $m_0 \leftarrow \widetilde{D}(n_0)$

REST OF THE PROTOCOL
 CHECKS $m_0 = \widetilde{w}_0(n_0)$

w_i already with $\widetilde{D}$

Put $q_i(0), q_i(1)$ in place of $\widetilde{w}_{i+1}(b), \widetilde{w}_{i+1}(c)$ to and Σ -check privacy
 $m_0 = \widetilde{w}_0(n_0)$

$q_i^* \leftarrow \mathbb{F}$
 $n_{i+1} = \ell(n_i^*)$
 $m_{i+1} = q_i(n_{i+1})$

Final check
 $m_d = \widetilde{w}_d(a_d) = \widetilde{w}_d(n_d)$

- $S_i = \# \text{ gates at level } i, k_i = \log_2 S_i$ (size to represent)
- $\widetilde{D}$ is the function of all the o/p values, there can be multiple outputs, $D = w_0$
- $\widetilde{D}$ makes the $\text{MLE}(D) = \widetilde{D}$
- $\widetilde{D}$ also makes the $\text{MLE}(\text{add}), \text{MLE}(\text{mult})$ since he knows the circuit already.
- But not $\text{MLE}(w_{i+1})$ because he doesn't know w_{i+1} , hence $\widetilde{D}$ does not know $f_{a_i}^{(1)}$
- $\widetilde{D}$ does not need to know $w_{i+1}()$ until the last round of Σ -check.
- D, w_i are functions mapping o/p gate labels to o/p gate values.
- No w_i are sent by the $\widetilde{D}$ to P .

m_{i+1} is sent and back to prover.

$\widetilde{w}_d() = \widetilde{w}_d()$ is the MLE of all inputs, which is known to both P & $\widetilde{D}$.
 $n_i: \{0, 1\}^d \rightarrow \mathbb{F}$

5 Questions

1. Complexity of prover and verifier?

ANS.

Communication	Rounds	$\mathcal{V}$ time	$\mathcal{P}$ time	Soundness error
$d \cdot \text{polylog}(S)$ field elements	$d \cdot \text{polylog}(S)$	$O(n + d \cdot \text{polylog}(S))$	$\text{poly}(S)$	$O(d \log(S)/ \mathbb{F})$

Table 4.4: Costs of the original GKR protocol [GKR08] when applied to any log-space uniform layered arithmetic circuit $\mathcal{C}$ of size S and depth d over n variables defined over field $\mathbb{F}$. Section 4.6.5 describes methods from [CMT12] for reducing $\mathcal{P}$'s runtime to $O(S \log S)$, and reducing the $\text{polylog}(S)$ terms in the remaining costs to $O(\log S)$. It is now known how to achieve prover runtime of $O(S)$ for arbitrary layered arithmetic circuits $\mathcal{C}$ (see Remark 4.5).

$\mathcal{V}$'s runtime. Observe that the polynomial $f_{r_i}^{(i)}$ defined in Equation (4.18) is a $(2k_{i+1})$ -variate polynomial of degree at most 2 in each variable, and so the invocation of the sum-check protocol at iteration i requires $2k_{i+1}$ rounds, with three field elements transmitted per round. Thus, the total communication cost is $O(S_0 + d \log S)$ field elements where S_0 is the number of outputs of the circuit. The time cost to $\mathcal{V}$ is $O(n + d \log S + t + S_0)$, where t is the amount of time required for $\mathcal{V}$ to evaluate add_i and mult_i at a random input, for each layer i of $\mathcal{C}$. Here the n term is due to the time required to evaluate $\widetilde{W}_d(r_d)$, the S_0 term is the time required to read the vector of claimed outputs and evaluate the corresponding multilinear extension, the $d \log S$ term is the time required for $\mathcal{V}$ to send messages to $\mathcal{P}$ and process and check the messages from $\mathcal{P}$. For now, let us assume that t is a low-order cost and that $S_0 = 1$, so that $\mathcal{V}$ runs in total time $O(n + d \log S)$; we discuss this issue further in Section 4.6.6

Round complexity and communication cost. By direct inspection of the protocol description, there are $O(d \log S)$ rounds in the GKR protocol, and the total communication cost is $O(d \log S)$ field elements.

$\mathcal{P}$'s runtime. Analogously to the MATMULT protocol of Section 4.4 we give two increasingly sophisticated implementations of the prover when the sum-check protocol is applied to the polynomial $f_{r_i}^{(i)}$.

Method 1: $f_{r_i}^{(i)}$ is a v -variate polynomial for $v = 2k_{i+1}$. As in the analysis of Method 1 for implementing the prover in the matrix multiplication protocol from Section 4.4, $\mathcal{P}$ can compute the prescribed method in round j by evaluating $f_{r_i}^{(i)}$ at $3 \cdot 2^{v-j}$ points. It is not hard to see that $\mathcal{P}$ can evaluate $f_{r_i}^{(i)}$ at any point in $O(S_i + S_{i+1})$ time using techniques similar to Lemma 3.8. This yields a runtime for $\mathcal{P}$ of $O(2^v \cdot (S_i + S_{i+1}))$. Over all d layers of the circuit, $\mathcal{P}$'s runtime is bounded by $O(S^3)$.

2. What if the S_i not a power of 2, let's say $S_i=3$, then what should the k_i be? Would it be $\lceil \log_2 S_i \rceil$, or something else?

ANS. It is only needed to define the variables and find their MLE, and MLE require boolean field, so I think the ceiling function should be reasonable enough.

3. Why is the function f is $2k_{i+1}$ variate polynomial?

ANS. Because f is defined over 2 variables of the lower layer, each of which requires atleast $\log_2 S_{i+1} = k_{i+1}$ bits to be represented.

4. If the Verifier does not know the function $f_{r_i}^{(i)}$, how does he do the final check in the protocol, because the Σ -check protocol requires evaluation of f at a random point?

ANS. Using the reduction of multiple polynomial evaluations to one explained in prerequisites.

5. Why do we need the MLE of add and mult? Is it just so because we have changed the domain of W for $\widetilde{W}$ and add/mult and $\widetilde{W}$ have the same input variables? Is that the only reason or are there others?

ANS. The [Eq. 2] is valid using any extension of add_i and mult_i , that are multilinear in the first k_i variables, where $k_i = \log_2 S_i$. But the MLE of add and mult are critical to achieving a prover runtime that is nearly **linear** in S (circuit size, total number of gates),

6. The verifier does know the circuit and the inputs but does he know the intermediate gate values?

ANS. No, he gets a MLE of all the values at a particular level

7. Comparing the PLONK permutation check and the wiring predicate here, they certainly are doing the same job of storing info about the circuit wiring? What does PLONK permutation check doesn't do?

ANS. PLONK permutation check can't store the operation itself, it just says which variable is reused again.

8. In the beginning, the P and V agree on a log space uniform arithmetic circuit $\mathcal{C}$ of fan-in 2 over a finite field $\mathbb{F}$. What is a log space uniform algorithm?

ANS.

evaluation problem. In this problem, $\mathcal{V}$ and $\mathcal{P}$ first agree on a *log-space uniform arithmetic circuit* C of fan-in 2 over a finite field $\mathbb{F}$, and the goal is to compute the value of the output gate(s) of C . A log-space uniform circuit C is one that possesses a succinct implicit description, in the sense that there is a logarithmic-space algorithm that takes as input the label of a gate a of C , and is capable of determining all relevant information about that gate. That is, the algorithm can output the labels of all of a 's neighbors, and is capable of determining if a is an addition gate or a multiplication gate.

9. How is the function $D(r_0)$ sent by the prover at the beginning claimed to be equal to the $W(r_0)$?

ANS. The verifier finds the MLE of D , $\tilde{D}$, and picks a random $r_0 \in F^{k_0}$, and evaluate $\tilde{D}(r_0)$ in time $O(S_0)$. Schwartz Zippel lemma if $\tilde{D}(r_0) = \tilde{W}(r_0)$, then it safe to assume that these are the same polynomial. Remember that the MLEs are distance amplifying algorithms, if two function differ even slightly, they would differ almost everywhere in their MLE.

10. who finds the $add(r, b, c)$ and $mult(r, b, c)$, the prover or the verifier? **ANS.** Verifier, because the circuit is known to both so why would prover do so.